
Machine Learning

Gaussian Mixture Models

Discriminative vs Generative Models

- Discriminative: Just learn a decision boundary between your sets.

Support Vector Machines


- Generative: Learn enough about your sets to be able to make new examples that would be set members

Gaussian Mixture Models

The Generative Model POV

- Assume the data was generated from a process we can model as a probability distribution
- Learn that probability distribution
- Once learned, use the probability distribution to
 - “Make” new examples
 - Classify data we haven’t seen before.

Non-parametric distribution not feasible




- Let's probabilistically model ML student heights.
- Ruler has 200 marks (100 to 300 cm)
- How many probabilities to learn?
- How many students in the class?
- What if the ruler is continuous?

Learning a Parametric Distribution


- Pick a parametric model (e.g. Gaussian)
- Learn just a few parameter values

$p(x | \Theta) \equiv$ prob. of x , given parameters Θ
of a model, M



Using Generative Models for Classification

Gaussians whose means and variances were learned from data



Answer: the class that calls him most probable.


Learning a Gaussian Distribution

$p(x | \Theta) \equiv$ prob. of x , given parameters Θ
of a model, M

$\Theta \equiv \{\mu, \sigma\}$ ← **The parameters we must learn**

$$M \equiv \frac{1}{(2\pi)^{1/2} \sigma} e^{\frac{-(x-\mu)^2}{2\sigma^2}}$$

↑
**The “normal” Gaussian distribution,
often denoted N , for “normal”**



Goal: Find the best Gaussian

- Hypothesis space is Gaussian distributions.
- Find parameters Θ^* that maximize the prob. of observing data $X \equiv \{x_1, \dots, x_n\}$

$$\Theta^* = \underset{\text{argmax } \Theta}{p(X | \Theta)}$$

where each $\Theta \equiv \{\mu, \sigma\}$

Some math

$$\Theta^* = \underset{\text{argmax } \Theta}{p(X | \Theta)}, \text{ where each } \Theta \equiv \{\mu, \sigma\}$$

$$p(X | \Theta) = \prod_{i=1}^n p(x_i | \Theta)$$

...if can we assume all x_i are i.i.d.

Numbers getting smaller

$$p(X | \Theta) = \prod_{i=1}^n p(x_i | \Theta)$$

What happens as n grows? Problem?

We get underflow if n is, say, 500

$$p(X | \Theta) \propto \sum_{i=1}^n \log(p(x_i | \Theta)) \text{ solves underflow.}$$

Remember what we're maximizing

$$\Theta^* \equiv \underset{\text{argmax } \Theta}{p(X | \Theta)} = \sum_{i=1}^n \underset{\text{argmax } \Theta}{\log(p(x_i | \Theta))}$$

fitting the Gaussian into this...

$$\log(p(x | \Theta)) = \log \left(\frac{e^{\frac{-(x-\mu)^2}{2\sigma^2}}}{(2\pi)^{1/2} \sigma} \right)$$

Some math gets you...

$$\log \left(\frac{e^{\frac{-(x-\mu)^2}{2\sigma^2}}}{(2\pi)^{1/2} \sigma} \right) = \log \left(e^{\frac{-(x-\mu)^2}{2\sigma^2}} \right) - \log((2\pi)^{1/2} \sigma)$$
$$= \boxed{\frac{-(x-\mu)^2}{2\sigma^2}} - \log \sigma - \log(2\pi)^{1/2}$$

Plug back into equation from slide 11

..which gives us

$$\Theta^* \equiv \underset{\text{argmax } \Theta}{p(X | \Theta)}$$

$$= \underset{\text{argmax } \Theta}{\sum_{i=1}^n \log(p(x_i | \Theta))}$$

$$= \underset{\text{argmax } \Theta}{\sum_{i=1}^n \left(\frac{-(x_i - \mu)^2}{2\sigma^2} - \log \sigma \right)}$$

Maximizing Log-likelihood

- To find best parameters, take the partial derivative with respect to parameters $\{\sigma, \mu\}$ and set to 0.

$$\Theta^* = \sum_{i=1}^n \left(\frac{-(x_i - \mu)^2}{2\sigma^2} - \log \sigma \right)$$

argmax Θ


- The result is a closed-form solution

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i \qquad \sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

What if...

- ...the data distribution can't be well represented by a single Gaussian?
- Can we model more complex distributions using multiple Gaussians?

Gaussian Mixture Model (GMM)



Model the distribution as a mix of Gaussians

$$P(x) = \sum_{j=1}^K P(z_j) P(x | z_j)$$

x is the observed value z_j is the j th Gaussian

What are we optimizing?

$$P(x) = \sum_{j=1}^K P(z_j) P(x | z_j)$$

Notating $P(z_j)$ as weight w_j and using the Normal (a.k.a. Gaussian) distribution $N(\mu_j, \sigma_j^2)$ gives us...

$$= \sum_{j=1}^K w_j N(x | \mu_j, \sigma_j^2) \quad \text{such that } 1 = \sum_{j=1}^K w_j$$

This gives 3 variables per Gaussian to optimize:

$$w_j, \mu_j, \sigma_j$$

Bad news: No closed form solution.

$$\begin{aligned}\Theta^* &\equiv \underset{\Theta}{\operatorname{argmax}} p(X | \Theta) = \underset{\Theta}{\operatorname{argmax}} \sum_{i=1}^n \log(p(x_i | \Theta)) \\ &= \underset{\Theta}{\operatorname{argmax}} \sum_{i=1}^n \log \left(\sum_{j=1}^K w_j p(x_i | N(\mu_j, \sigma_j^2)) \right)\end{aligned}$$

Expectation Maximization (EM)

- Solution: The EM algorithm
- EM updates model parameters iteratively.
- After each iteration, the likelihood the model would generate the observed data increases (or at least it doesn't decrease).
- EM algorithm always converges to a local optimum.

EM Algorithm Summary

- Initialize the parameters
- E step: calculate the likelihood a model with these parameters generated the data
- M step: Update parameters to increase the likelihood from E step
- Repeat E & M steps until convergence to a local optimum.

EM for GMM - Initialization

- Choose the number of Gaussian components K
K should be much less than the number of data points to avoid overfitting.
- (Randomly) select parameters for each Gaussian j : w_j, μ_j, σ_j

...such that $1 = \sum_{j=1}^K w_j$

EM for GMM – Expectation step

The responsibility $\gamma_{j,n}$ of Gaussian j for observation x_n is defined as...

$$\begin{aligned}\gamma_{j,n} &\equiv p(z_j | x_n) = \frac{p(x_n | z_j)p(z_j)}{p(x_n)} \\ &= \frac{p(x_n | z_j)p(z_j)}{\sum_{k=1}^K p(z_k)p(x_n | z_k)} = \frac{w_j N(x_n | \mu_j, \sigma_j^2)}{\sum_{k=1}^K w_k N(x_n | \mu_k, \sigma_k^2)}\end{aligned}$$

EM for GMM – Expectation step

Define the responsibility Γ_j of Gaussian j for all the observed data as...

$$\Gamma_j \equiv \sum_{n=1}^N \gamma_{j,n}$$

You can think of this as the proportion of the data explained by Gaussian j .

EM for GMM – Maximization step

Update our parameters as follows...

$$\begin{aligned}\text{new } w_j &= \frac{\Gamma_j}{N} \\ \text{new } \mu_j &= \frac{\sum_{i=1}^N \gamma_{j,i} x_i}{\Gamma_j} \\ \text{new } \sigma_j^2 &= \frac{\sum_{i=1}^N \gamma_{j,i} (x_i - \mu_j)^2}{\Gamma_j}\end{aligned}$$

Why does this work?

- We need to prove that, as our model parameters are adjusted, likelihood of the data never goes down (monotonically non-decreasing)
- This is the part where I point you to the textbook

What if...

- ...our data isn't just scalars, but each data point has multiple dimensions?
- Can we generalize to multiple dimensions?
- We need to define a covariance matrix.

Covariance Matrix

Given d-dimensional random variable vector $\vec{\mathbf{X}} = [X_1, \dots, X_d]$
the covariance matrix denoted Σ (confusing, eh?) is defined as...


$$\Sigma \equiv \begin{bmatrix} \mathbf{E}[(X_1 - \mu_1)(X_1 - \mu_1)] & \mathbf{E}[(X_1 - \mu_1)(X_2 - \mu_2)] & \dots & \mathbf{E}[(X_1 - \mu_1)(X_d - \mu_d)] \\ \mathbf{E}[(X_2 - \mu_2)(X_1 - \mu_1)] & \mathbf{E}[(X_2 - \mu_2)(X_2 - \mu_2)] & \dots & \mathbf{E}[(X_2 - \mu_2)(X_d - \mu_d)] \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{E}[(X_d - \mu_d)(X_1 - \mu_1)] & \mathbf{E}[(X_d - \mu_d)(X_2 - \mu_2)] & \dots & \mathbf{E}[(X_d - \mu_d)(X_d - \mu_d)] \end{bmatrix}$$

This is a generalization of one-dimensional variance for a scalar random variable X

$$\sigma^2 = \text{var}(X) = E[(X - \mu)^2]$$

Multivariate Gaussian Mixture

Second dimension




The d by d covariance matrix Σ describes the shape and orientation of an ellipse.

$$P(\vec{X}) = \sum_{j=1}^K w_j p(\vec{X} | N(\vec{\mu}, \Sigma_j))$$


Given d dimensions and K Gaussians, how many parameters?

Example: Initialization




(Illustration from Andrew Moore's tutorial slides on GMM)

After Iteration #1




(Illustration from Andrew Moore's tutorial slides on GMM)

After Iteration #2




(Illustration from Andrew Moore's tutorial slides on GMM)

After Iteration #3




(Illustration from Andrew Moore's tutorial slides on GMM)

After Iteration #4




(Illustration from Andrew Moore's tutorial slides on GMM)

After Iteration #5




(Illustration from Andrew Moore's tutorial slides on GMM)

After Iteration #6



(Illustration from Andrew Moore's tutorial slides on GMM)

After Iteration #20



(Illustration from Andrew Moore's tutorial slides on GMM)

GMM Remarks


- GMM is powerful: any density function can be arbitrarily-well approximated by a GMM with enough components.
- If the number of components K is too large, data will be overfitted.
 - Likelihood increases with K .
 - Extreme case: N Gaussians for N data points, with variances $\rightarrow 0$, then likelihood $\rightarrow \infty$.
- How to choose K ?
 - Use domain knowledge.
 - Validate through visualization.

GMM is a “soft” version of K-means

- Similarity
 - K needs to be specified.
 - Converges to some local optima.
 - Initialization matters final results.
 - One would want to try different initializations.
- Differences
 - GMM Assigns “soft” labels to instances.
 - GMM Considers variances in addition to means.

GMM for Classification

- Given training data with multiple classes...
 - 1) Model the training data for each class with a GMM
 - 2) Classify a new point by estimating the probability each class generated the point
 - 3) Pick the class with the highest probability as the label.



(illustration from
Leon Bottou's slides
on EM)

GMM for Regression

Given dataset $D = \{ \langle x_1, y_1 \rangle, \dots, \langle x_n, y_n \rangle \}$, where $y_i \in \mathcal{R}$ and x_i is a vector of d dimensions...

Learn a $d + 1$ dimensional GMM.

Then, compute $f(x) = \mathbf{E}[y | x]$

